cross-review-mcp

User Guide v0.5.2

Get expert advice from 5 AI models at the same time.

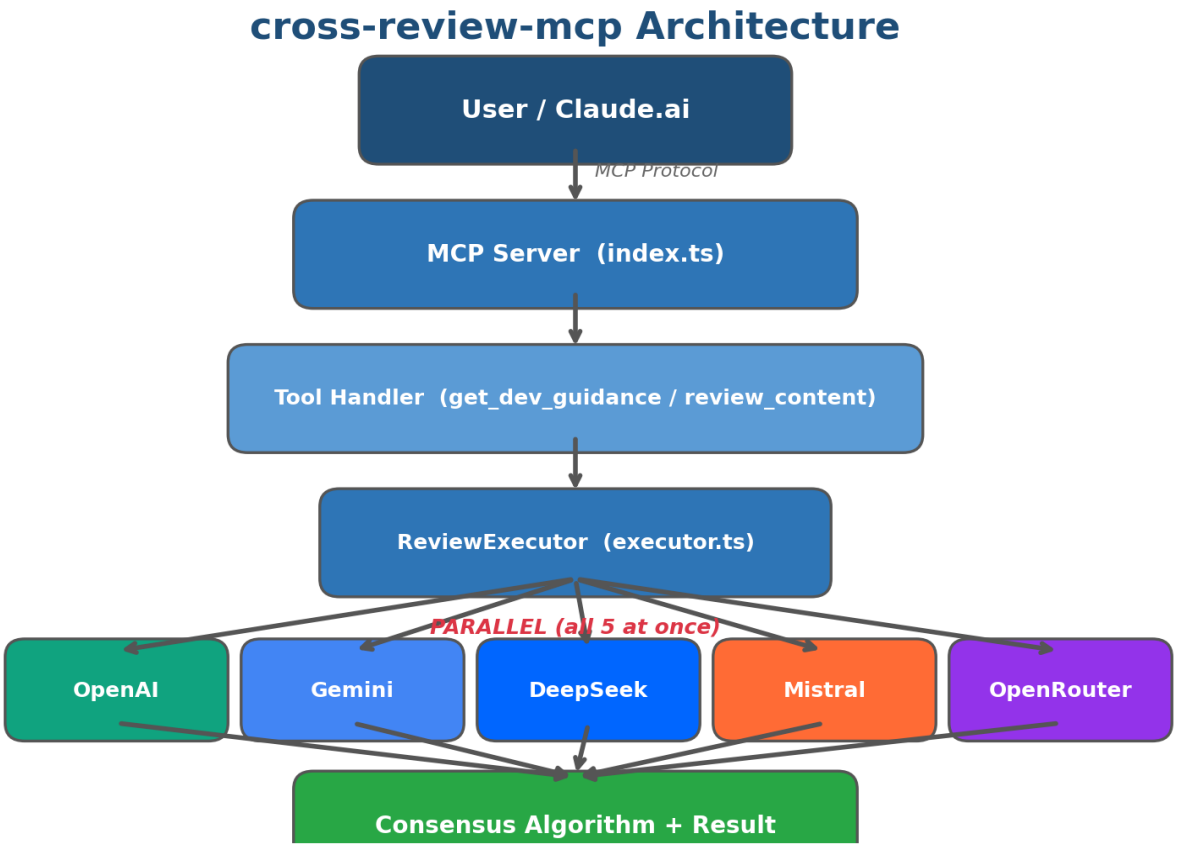

Figure 1: System architecture — your question goes to 5 models in parallel

- Ask 5 AI models in parallel — same speed as asking 1
- Consensus verdict shows what models agree on
- Divergent perspectives highlight alternative approaches
- Developer guidance tool for solving blockers instantly

What Is This?

Imagine asking 5 expert friends the same question at once. Some agree, some disagree — but you get **5 perspectives** instead of 1. That's cross-review-mcp.

Instead of asking ChatGPT OR Gemini OR Claude separately, you ask **all 5 AI models at the SAME TIME** and get a consensus verdict.

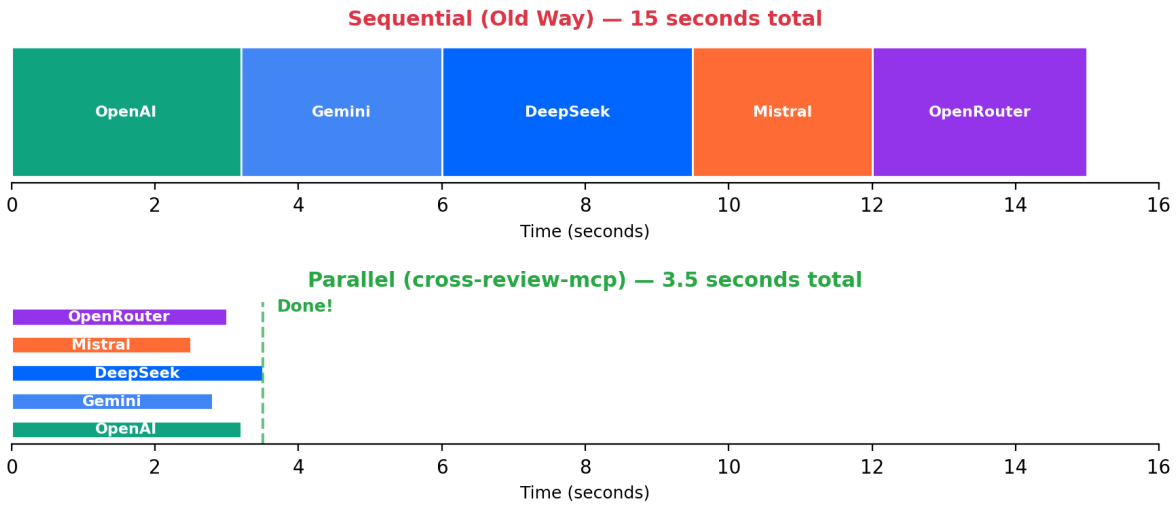

Figure 2: Sequential vs Parallel — same total work, 4x faster

Consensus in Action

When all 5 models analyze your problem, the consensus algorithm finds agreement and highlights where models diverge — because disagreement is often valuable.

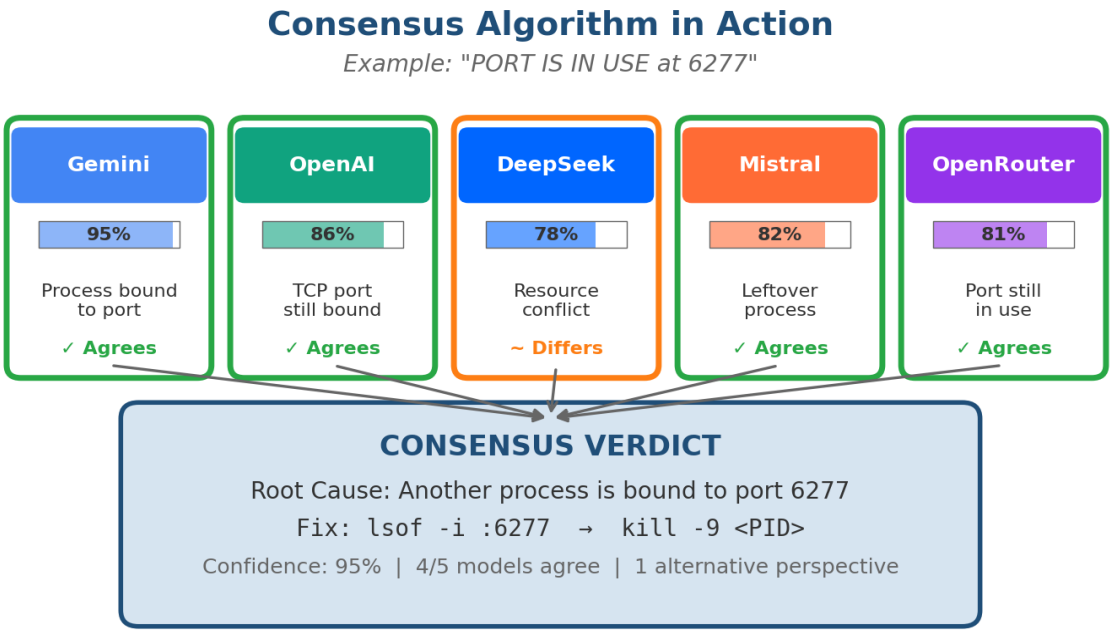

Figure 3: 4 of 5 models agree on root cause; 1 offers alternative angle

How It Works

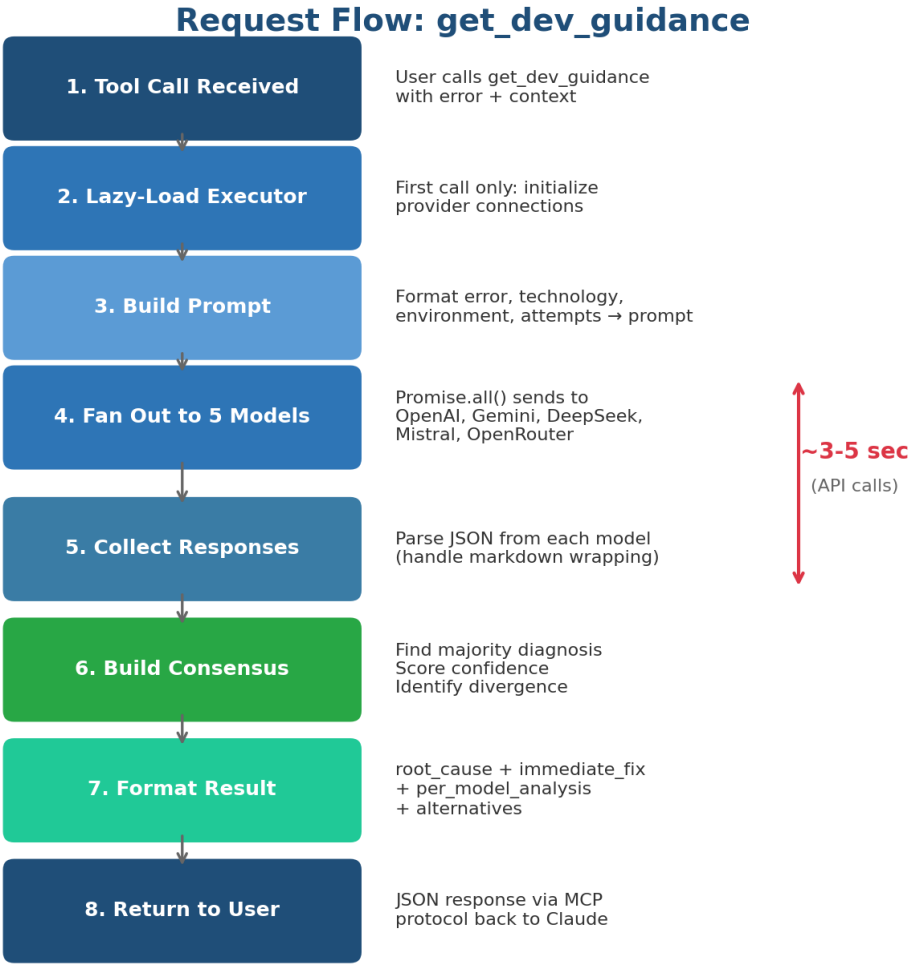

Figure 4: Complete request flow from question to answer

Installation (3 Steps)

Step 1: Install via npm

```
npm install -g cross-review-mcp
```

Step 2: Get API keys from: OpenAI, Gemini, DeepSeek, Mistral, OpenRouter

Step 3: Create a `.env` file:

```
OPENAI_API_KEY=sk-proj-... GEMINI_API_KEY=AIzaSy-... DEEPSEEK_API_KEY=sk-...
MISTRAL_API_KEY=... OPENROUTER_API_KEY=sk-or-v1-...
```

Don't have all 5? That's fine — use what you have.

Usage

CLI:

```
cross-review dev --error "PORT IS IN USE at 6277" \ --tech "MCP Inspector" --env "macOS" \
--tried "Killed processes"
```

Claude.ai: Connect via MCP, then ask naturally — Claude calls cross-review automatically.

Docker:

```
docker run \ -e OPENAI_API_KEY=sk-... \ -e GEMINI_API_KEY=AIza... \ cross-review-mcp
```

Cost & Performance

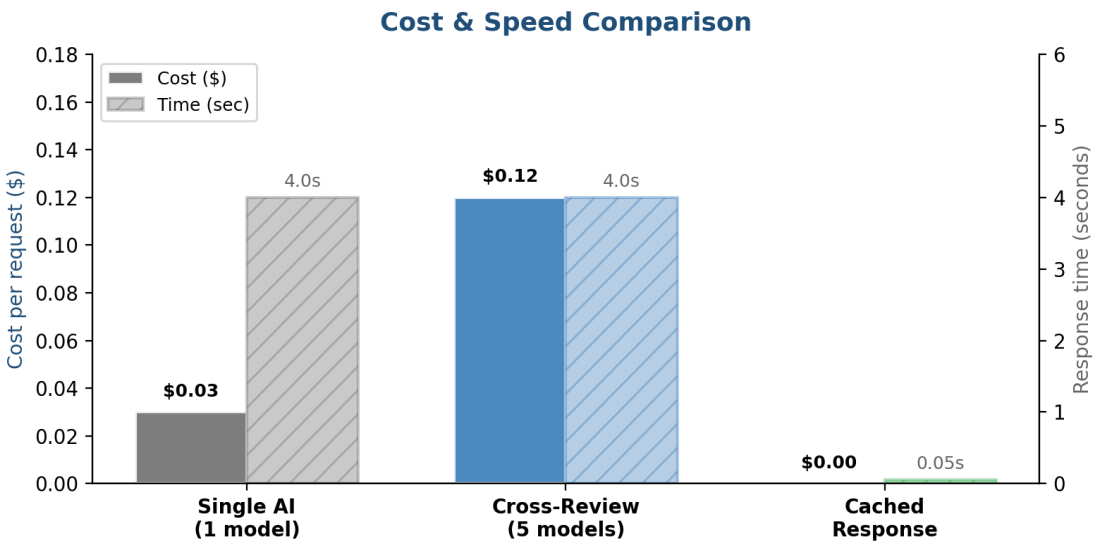

Figure 5: Cost and speed — 5 models cost ~$0.12 but take same time as 1

Operation	Time	Cost	Models
Single AI call	3-5 sec	$0.01-0.05	1
Cross-Review (5 models)	3-5 sec	$0.05-0.25	5
Cached response	<0.1 sec	$0.00	0

Model Capabilities

Each AI model has different strengths. The radar chart below shows how models performed on a real problem (PORT IS IN USE), across 5 capability dimensions.

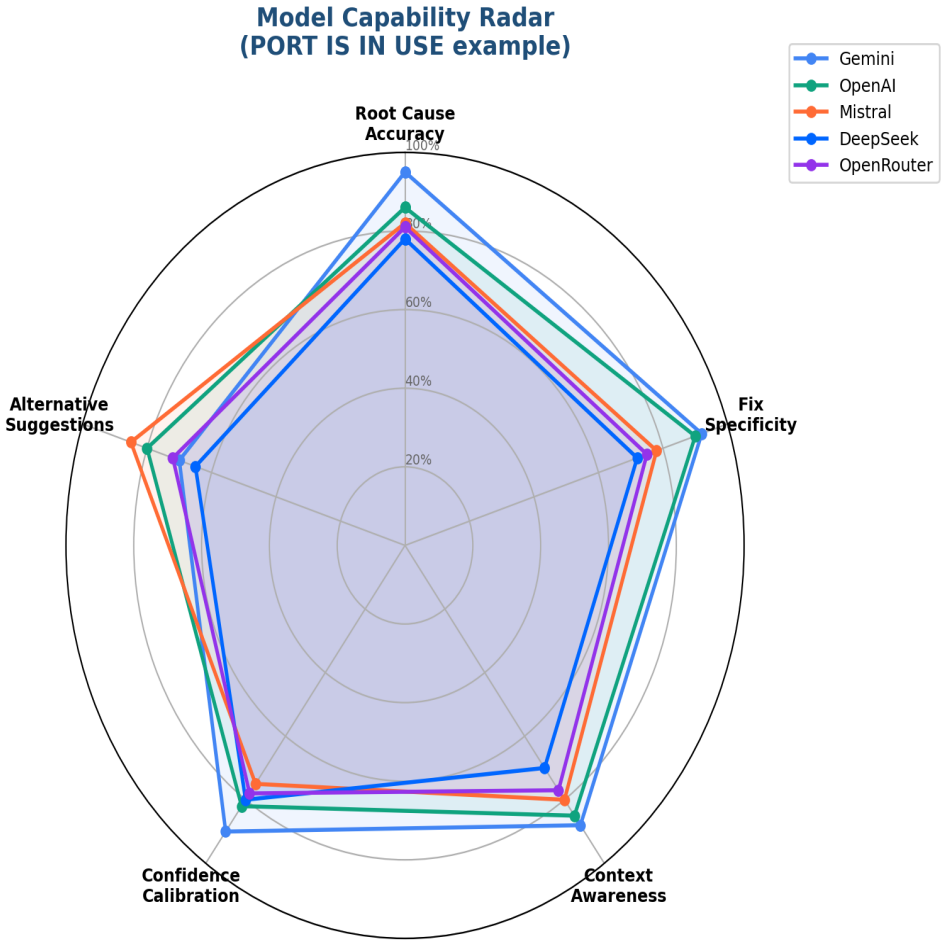

Figure 6: Model capability radar — Gemini excels at root cause; OpenAI at alternatives

Troubleshooting

"Error: Missing API key" — Create .env file with your keys. See Installation above.

"No model responses" — Verify API keys are correct and not expired. Check network.

"Answer isn't helpful" — Add more context: technology, environment, what you've tried.

"Timeout" — Some providers may be slow. The system waits for all 5 by default.